

A Maximum-Entropy Derivation of Universal Structural Selection

From Constraint Defects to Exponential Structural Measures

Christof Krieg

MAAT Project – Independent Research

<https://maat-research.com>

2026

Abstract

The MAAT/Structural Selection series has used structural functionals to rank field configurations, cosmological histories, computational instances, and string-landscape candidates. This paper gives the formal reason why this class of functionals has a universal exponential form.

Let $\mathcal{X}_{\text{adm}}$ be a space of dynamically admissible configurations with reference measure μ_0 , and let $d_a[X] \geq 0$ be measurable structural defects encoding consistency, balance, activity, connectivity, robustness, or domain-specific closure constraints. Among all probability densities $p[X]$ over μ_0 with fixed mean defects $\langle d_a \rangle = \bar{d}_a$, the unique maximum-relative-entropy solution is

$$p_\lambda[X] = \frac{1}{Z(\lambda)} \exp \left[- \sum_a \lambda_a d_a[X] \right].$$

Thus the linear structural functional

$$F_{\text{struct}}[X] = \sum_a \lambda_a d_a[X]$$

is not an arbitrary score: it is the Gibbs–Jaynes rate functional generated by structural constraints.

The support form used in later MAAT papers,

$$F_{\text{MAAT}}[X] = - \sum_a \lambda_a \log(\epsilon + \Gamma_a[X]), \quad \Gamma_a[X] = \frac{1}{1 + d_a[X]},$$

is shown to be a logarithmic support representation of the same selection logic. For small defects it reduces locally to a linear defect cost, while globally it gives a bounded support interpretation in which each sector contributes a multiplicative viability factor.

The theorem proves the form of the selection measure conditional on the choice of defects, reference measure, and constraints. It does not prove that a particular defect basis is fundamental. That stronger universality claim must be tested by recurrence across domains and by outperforming simpler alternatives such as energy-only, entropy-only, stability-only, random, anthropic, or complexity-only rankings.

1 Purpose and Status

The Structural Selection programme has accumulated evidence that energy alone does not fully characterize which configurations are coherent, stable, or physically relevant. Examples include kink selection in nonlinear fields, fixed-energy structural splitting, cosmological structural-stress observables, SAT hardness via defect covariance, and string-landscape rankings that differ from energy ordering.

The present paper is not another numerical benchmark. It is the formal closure step:

If realized configurations are selected under finite information about structural defects, then the least-biased selection measure is exponential in an additive structural functional.

This proves a conditional theorem. The theorem derives the form of the probability measure once a reference measure and structural constraints are specified. It does not derive the microscopic origin of every defect sector.

2 Admissible Configurations and Reference Measure

Let $\mathcal{X}_{\text{adm}}$ be a measurable space of admissible configurations. Depending on the application, $X \in \mathcal{X}_{\text{adm}}$ may be:

- a field configuration satisfying coarse dynamical admissibility;
- a cosmological history;
- a SAT or constraint-system instance;
- a candidate compactification or string background;
- a reduced effective trajectory in an AI or control system.

Let μ_0 be a reference measure on $\mathcal{X}_{\text{adm}}$. It represents the baseline ensemble before structural selection is imposed. Examples include a lattice field measure, an uncontrolled simulation ensemble, a flux-scan prior, an instance-generation distribution, or a coarse-grained empirical ensemble.

Let

$$d_a : \mathcal{X}_{\text{adm}} \rightarrow [0, \infty)$$

be structural defect observables. The index a labels sectors such as equation consistency, constraint balance, activity, connectivity, robustness, or domain-specific diagnostics.

3 Maximum-Entropy Structural Selection Theorem

Theorem. Assume that the moment-generating partition function

$$Z(\lambda) = \int_{\mathcal{X}_{\text{adm}}} \exp \left[- \sum_a \lambda_a d_a[X] \right] d\mu_0[X]$$

is finite for a set of Lagrange multipliers λ_a . Among all probability densities $p[X]$ with respect to μ_0 satisfying

$$\int_{\mathcal{X}_{\text{adm}}} p[X] d\mu_0[X] = 1, \quad \int_{\mathcal{X}_{\text{adm}}} p[X] d_a[X] d\mu_0[X] = \bar{d}_a,$$

the density maximizing relative entropy

$$S[p|\mu_0] = - \int_{\mathcal{X}_{\text{adm}}} p[X] \log p[X] d\mu_0[X]$$

is

$$\boxed{p_\lambda[X] = \frac{1}{Z(\lambda)} \exp \left[- \sum_a \lambda_a d_a[X] \right]}. \quad (1)$$

Proof. Introduce Lagrange multipliers α and λ_a and vary

$$\mathcal{L}[p] = - \int p \log p \, d\mu_0 - \alpha \left(\int p \, d\mu_0 - 1 \right) - \sum_a \lambda_a \left(\int p \, d_a \, d\mu_0 - \bar{d}_a \right).$$

Stationarity with respect to p gives

$$-(\log p[X] + 1) - \alpha - \sum_a \lambda_a d_a[X] = 0.$$

Therefore

$$p[X] = \exp[-1 - \alpha] \exp \left[- \sum_a \lambda_a d_a[X] \right].$$

Normalization fixes $\exp[1 + \alpha] = Z(\lambda)$, yielding Eq. (1). Strict concavity of entropy on the convex set of admissible densities gives uniqueness when the constraints are feasible. $\square$

4 Structural Functional as Rate Functional

Equation (1) defines the linear structural functional

$$F_{\text{struct}}[X] = \sum_a \lambda_a d_a[X]. \quad (2)$$

Then

$$p_\lambda[X] = \frac{1}{Z(\lambda)} e^{-F_{\text{struct}}[X]}. \quad (3)$$

This is the direct analogue of the Gibbs distribution in statistical mechanics. Energy is replaced by structural defect cost, and thermodynamic intensive variables are replaced by structural selection pressures λ_a .

Large-deviation theory gives the same interpretation in asymptotic ensembles: configurations concentrate near minima of the effective rate functional. Thus structural selection can be read either as a MaxEnt inference principle or as a large-deviation concentration principle.

5 Support-Form Representation

Later MAAT papers often use structural supports

$$\Gamma_a[X] = \frac{1}{1 + d_a[X]}, \quad 0 < \Gamma_a \leq 1. \quad (4)$$

The support-form functional is

$$F_{\text{MAAT}}[X] = - \sum_a \lambda_a \log(\epsilon + \Gamma_a[X]). \quad (5)$$

It gives

$$e^{-F_{\text{MAAT}}[X]} = \prod_a (\epsilon + \Gamma_a[X])^{\lambda_a}. \quad (6)$$

Thus each structural sector contributes a multiplicative support factor.

For small defects and $\epsilon \ll 1$,

$$\log \Gamma_a = -\log(1 + d_a) = -d_a + O(d_a^2),$$

so

$$-\lambda_a \log(\epsilon + \Gamma_a) \simeq \lambda_a d_a \quad \text{up to constants and higher-order terms.} \quad (7)$$

Therefore the support form is locally equivalent to the linear MaxEnt functional, while globally it provides a bounded viability interpretation.

6 Weight Closure

The MaxEnt theorem proves the exponential family form but does not by itself fix the numerical values of λ_a . The weights are Lagrange multipliers conjugate to the chosen constraints:

$$\bar{d}_a = -\frac{\partial \log Z}{\partial \lambda_a}.$$

In the response-closure approximation used in later papers,

$$\lambda \approx \left(C + \eta \frac{\text{tr}(C)}{A} \mathbb{I} \right)^{-1} (\langle d \rangle_{\mu_0} - \langle d \rangle_*), \quad (8)$$

where $C = \text{Cov}_{\mu_0}(d)$ is the defect covariance matrix, A is the number of sectors, and $\langle d \rangle_*$ is a target low-defect or safe subensemble. This makes λ_a response pressures rather than arbitrary constants.

MaxEnt itself does not require every multiplier to be positive. The sign of λ_a is fixed by the imposed target moments and by the orientation of the corresponding defect observable. For structural-selection applications one normally chooses defect conventions in which larger d_a means larger obstruction, so viable selection pressures are expected to be non-negative. Negative or unstable weights would therefore be interpreted as evidence that the defect definition, target ensemble, or reference measure has been chosen inconsistently.

7 Universal Structural Selection Claim

The theorem supports the following conditional universality statement:

Any system that selects realized configurations under limited information about additive structural constraints has a least-biased selection measure exponential in a structural functional.

The stronger claim that MAAT is universal requires two additional steps:

1. identify recurring defect sectors across domains;
2. show that the resulting structural functional predicts selected configurations better than simpler alternatives.

Thus the formal proof and the empirical universality programme are distinct. The proof gives the shape of the measure. The benchmarks test whether the chosen defects are useful, recurrent, and nontrivial.

8 Evidence Ladder Across Domains

Table 1: Examples of recurring structural-selection tests in the series.

Domain	Selection object	Evidence type
Nonlinear fields	Kinks, domain walls, fixed-energy field states	Structure ranking separates configurations beyond energy alone.
Cosmology	FLRW histories, CCI, projection observables, growth response	Structural observables define testable stress, growth, and metric-response channels.
SAT / constraints	SAT hardness and boundary regimes	Defect covariance and robustness closure identify hardness/constraint geometry.
Effective constants	SM-like RG bridge	Low-defect basins correlate with observed effective constants without exact derivation.
String landscape	KKLT vacua and transition paths	Structural vacuum ranking and structural path flow differ from energy-only ordering.
AI / safety	Robustness and safety-boundary ensembles	Response weights become dominated by constraint-respect under safety targets.

This table is not a proof of universality. It is the empirical ladder required after the MaxEnt theorem.

9 Open Universality Tests

Several domains remain especially important because they would test whether the same structural-selection form survives outside the systems used to develop it:

- **Quantum measurement:** pointer-state selection should be tested against a structural support functional over decohering branches, not merely assumed by analogy.
- **Independent computational datasets:** SAT and constraint-system closures should be tested on generators and solver families not used in the original calibration.
- **Cosmological likelihoods:** structural observables should be embedded in standard likelihood pipelines rather than only compared through diagnostic proxies.
- **String landscape ensembles:** vacuum and path selection should be tested on ensembles beyond reduced KKLT-style scans.

These tests are deliberately phrased as possible failure modes. The MaxEnt theorem protects the form of the measure, but not the correctness of any proposed defect basis.

10 Falsification Criteria

A universal structural-selection principle must be falsifiable. The following tests are therefore essential:

- compare F_{MAAT} against energy-only ranking;

- compare against entropy-only or activity-only scores;
- compare against stability-only ranking;
- compare against random or shuffled defect rankings;
- compare against anthropic or hand-selected filters where applicable;
- test out-of-sample defect weights and response closures;
- test whether domain-specific defects add predictive power beyond the shared core sectors.

The programme fails if F_{MAAT} repeatedly collapses to a trivial proxy for energy, entropy, or a single stability variable. It succeeds only if the same exponential structural form continues to add predictive information across independent domains.

11 What Has Been Proved and What Has Not

Table 2: Scope of the present derivation.

Established here	Not established here
MaxEnt implies $P[X] \propto e^{-F_{\text{struct}}[X]}$ under fixed defect means.	The microscopic origin of every defect sector.
The linear defect functional is the Gibbs–Jaynes structural cost.	That the MAAT defect basis is unique.
The support form is a logarithmic/multiplicative support representation.	That all physical systems must use exactly the same sector set.
Response closure interprets weights as covariance response pressures.	Full first-principles derivation of λ_a from quantum gravity, renormalization, or entropy production.

12 Reproducibility and Repository

This paper is primarily formal and introduces no new numerical solver. The theorem can be checked directly from the variational calculation in Section 3. The empirical evidence ladder cited here is reproducible through the public Structural Selection repository:

<https://github.com/Chris4081/structural-selection-principle>

Relevant existing benchmark folders include:

- the cosmological CCI folders for structural-stress diagnostics;
- the growth, projection, and metric-response folders for observable cosmology benchmarks;
- the string-landscape ranking and path-selection folders for vacuum and transition-flow tests;
- the void-environment folder for environment-dependent metric-response benchmarks.

All repository CSV, JSON, and PNG files are derived reproducibility artifacts of the corresponding scripts and should be cited together with the papers and original external data sources used in each benchmark.

13 Conclusion

This paper derives the formal core of universal structural selection. If admissible configurations are selected under finite information about structural defects, the least-biased probability measure is necessarily exponential in an additive structural functional:

$$P[X] \propto e^{-F_{\text{struct}}[X]}.$$

The MAAT support form is a bounded logarithmic representation of the same selection logic.

The result does not prove that MAAT is a final theory of nature. It proves the mathematical form that any MaxEnt structural-selection measure must take. The universality claim must then be tested by recurrence, out-of-sample prediction, and falsification against simpler competing rankings.

References

- [1] E. T. Jaynes, “Information Theory and Statistical Mechanics,” *Physical Review* **106**, 620–630 (1957). DOI: [10.1103/PhysRev.106.620](https://doi.org/10.1103/PhysRev.106.620).
- [2] E. T. Jaynes, “Information Theory and Statistical Mechanics. II,” *Physical Review* **108**, 171–190 (1957). DOI: [10.1103/PhysRev.108.171](https://doi.org/10.1103/PhysRev.108.171).
- [3] S. Kullback and R. A. Leibler, “On Information and Sufficiency,” *Annals of Mathematical Statistics* **22**, 79–86 (1951). DOI: [10.1214/aoms/1177729694](https://doi.org/10.1214/aoms/1177729694).
- [4] T. M. Cover and J. A. Thomas, *Elements of Information Theory*, 2nd ed., Wiley (2006).
- [5] R. S. Ellis, *Entropy, Large Deviations, and Statistical Mechanics*, Springer (1985).
- [6] H. Touchette, “The large deviation approach to statistical mechanics,” *Physics Reports* **478**, 1–69 (2009). DOI: [10.1016/j.physrep.2009.05.002](https://doi.org/10.1016/j.physrep.2009.05.002).
- [7] C. Krieg, “From Structural Diagnostics to Field-Theoretic Selection Functionals,” preprint (2026).
- [8] C. Krieg, “Response-Based Derivation of MAAT Structural Weights: From Covariance Geometry to Selection Pressure,” preprint (2026).